

# Altera's Partial Reconfiguration Flow

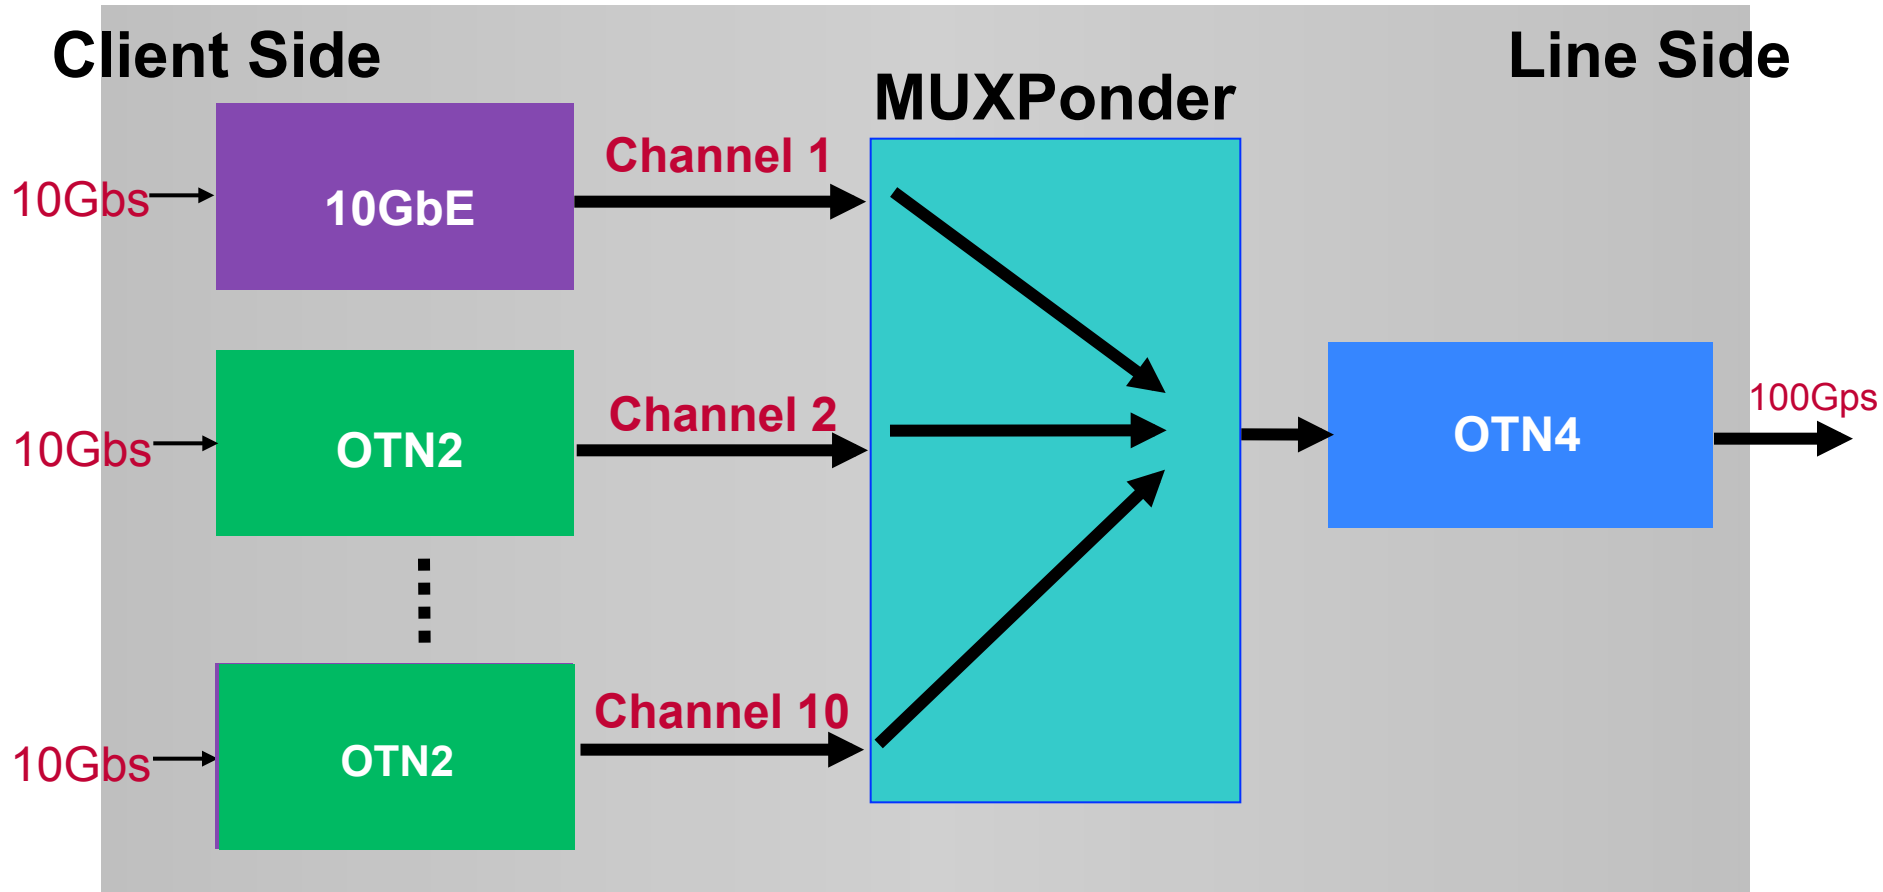
Mark Bourgeault

June 23, 2011

# Outline




- High Level Overview
- Hardware Support
- Reconfigurable Projects
- Software Implementation Details
- Upcoming Features & Summary

# Example System: 10\*10Gbps→OTN4 Muxponder





# Altera Partial Reconfiguration Strategy

## ■ Software flow is key

-  Enter *design intent*, automate low-level details
-  Build on existing incremental design & floorplanning tools
-  Intuitive simulation flow, *including effect of reconfiguration*

## ■ Partial reconfiguration can be controlled by soft logic, or an external device


-  Load partial programming files while device operating
-  Support Configuration Via Protocol (CvP, e.g. PCIe)

## ■ Targets **multi-modal** applications

-  Coarse-grained reconfiguration

# Partial Reconfiguration Sequence






- Bring-up the chip with a known configuration
- Make decision to reprogram reconfigurable components
- Issue “PR request”
- Send “PR data”
  - Partial bitstream generated by Quartus
- Receive “PR ready”
- Rest of logic can now interact with reconfigured component




# Sample Design File Content

```
module reconfig_channel (clk, in, out);  
input clk, in;  
output [7:0] out;  
  
parameter VER = 2; // 1 to select 10GbE, 2 to select OTN2  
  
generate  
    case (VER)  
        1: gige m_gige (.clk(clk), .in(in), .out(out));  
        2: otn2 m_otn2 (.clk(clk), .in(in), .out(out));  
        default: gige m_gige(.clk(clk), .in(in), .out(out));  
    endcase  
endgenerate  
  
endmodule
```

# Stratix V Device Floorplan

-  I/O, Memory Interfaces, Clocks, etc...
-  PLL, JTAG, CRC, control circuitry, etc...
-  Transceivers, Hard IP Blocks
-  LAB, RAM, DSP
-  PLLs, Clocks



# Reconfiguration Modes Per Resources

Resources	Reconfiguration Type
Logic Block (LAB)	Partial Reconfiguration
Digital Signal Processing (DSP)	Partial Reconfiguration
Memory Block (RAM)	Partial Reconfiguration
Transceivers	Dynamic Reconfiguration (ALTGX_Reconfig)
PLL	Dynamic Reconfiguration (ALTPLL_Reconfig)
IO Blocks & Other Periphery	Altera Memory Interface IP



# Configuration RAM Bits

- FPGA is fixed HW with millions of 1-bit memory cells (CRAM)
  - Core bits are directly addressable by Frame (column) & Index (row)
  - Large devices have  $>10^7$  CRAMs
- By changing CRAM values, can reconfigure individual LABs, RAMs, DSPs, and *routing muxes*
  - 📁 Not all CRAM sequences are legal
  - 📁 Quartus will not provide details about CRAM → logical function

CRAM address space	Frame 1	Frame 2	...	...	Frame m	Frame m+1	Frame m+2	...	...	...	Frame n	Frame n+1	Frame n+2	...	...	...	Last Frame
Bit 1																	
Bit 2																	
⋮																	
⋮																	
⋮																	
Bit i																	
Bit i+1																	
⋮																	
⋮																	
Bit i+j-1																	
Bit i+j																	
Last Bit																	

Non-PR Region

PR Region

# Basic PR Operation

## ■ Use handshaking commands



- From user I/O or from core logic

## ■ Speed of operation

- Small overhead to engage PR
- Depends on size of region
- Depends on orientation of region


## ■ Sample configuration times

- ~5kLEs
- Region A ~ 2ms
- Region B ~ 10ms
- 2ms @ 400 MHz = 800,000 cycles



# Changing CRAM Bit State


- For each affected frame, control block generates
  - An AND instruction to set reconfigured bits to zero
  - An OR instruction to set them to one



# Implications of Temporary Frame


## ■ Must carefully consider electrical effects

- NMOS-based logic is very common
- Requires weak pull-up transistor to strengthen internal node when input signal is VDD
- Global “freeze” signal can’t be used during PR




# Implications of Temporary Frame

- 'AND' instruction may cause all CRAM bits of a routing driver to be zero
  - No electrical issues if 'x' is VDD at the time when the CRAM bit flips from a '1' to a '0'




# Implications of Temporary Frame

- 'AND' instruction may cause all CRAM bits of a routing driver to be zero
  - However, if 'x' is GND when the CRAM flips, voltage on 'y' can drift causing unwanted current



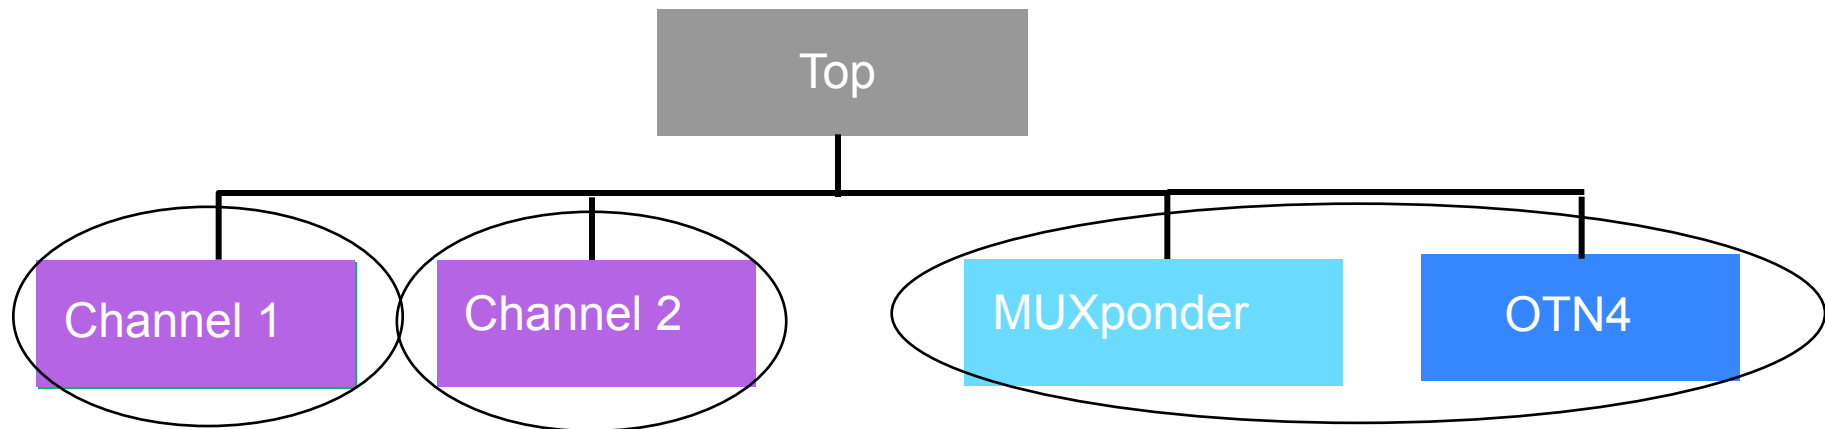
# Partial Reconfiguration Inputs & Outputs



***Multiple Compiles Needed!***

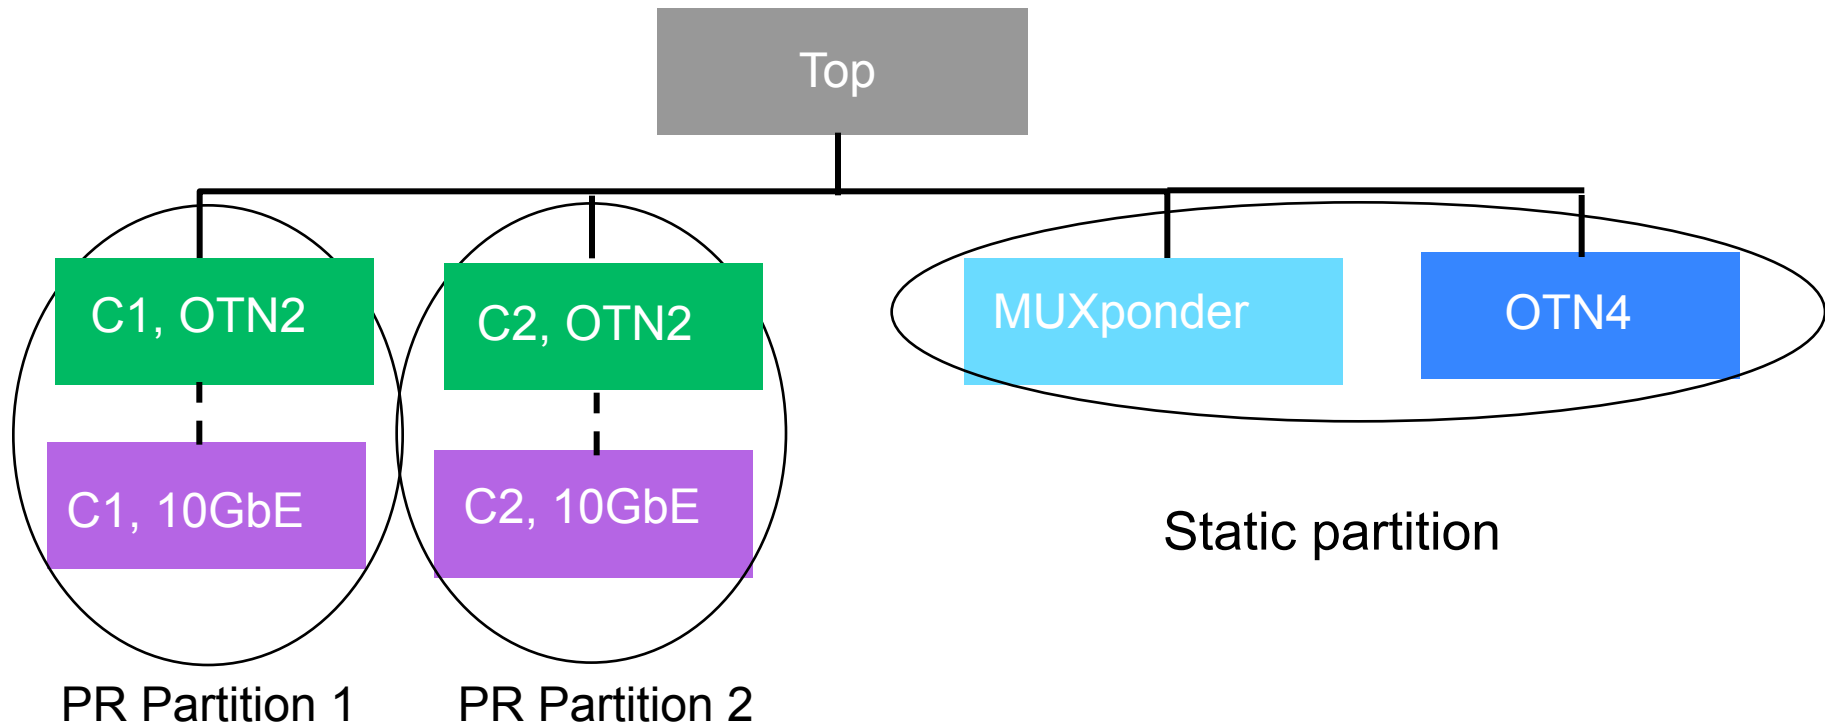
# Incremental Design Background

- Specify *partitions* in your design hierarchy
- Can independently recompile any partition
  - CAD optimizations across partitions prevented
  - Can preserve synthesis, placement and routing of unchanged partitions
  - Supports multiple processors for faster compile






# Reconfigurable Instances



■ Each design instance in a reconfigurable partition is called a “persona”


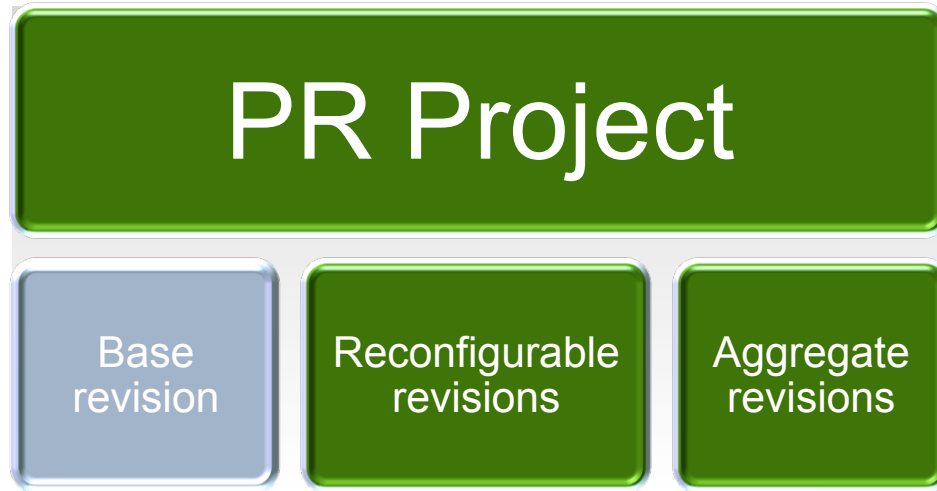
- Above design hierarchy has 4 personae (C1 x 2, C2 x 2)
- PR partitions must be assigned to fixed regions on the chip

# Partial Reconfiguration Project



- Revision = set of input files for a single compile
  - Can change parameters, timing constraints, fitter settings, etc...
- Revision types in a PR project
  - Base: Content that always exists on the chip (a.k.a. static compile)
  - Reconfigurable: Content that can be programmed via PR commands
  - Aggregate: User-defined groupings of reconfigurable revisions  
(consider 10 PR partitions with 6 personae each,  $10^6$  possible combos)
- No limit to the # of regions or # of personae / PR region


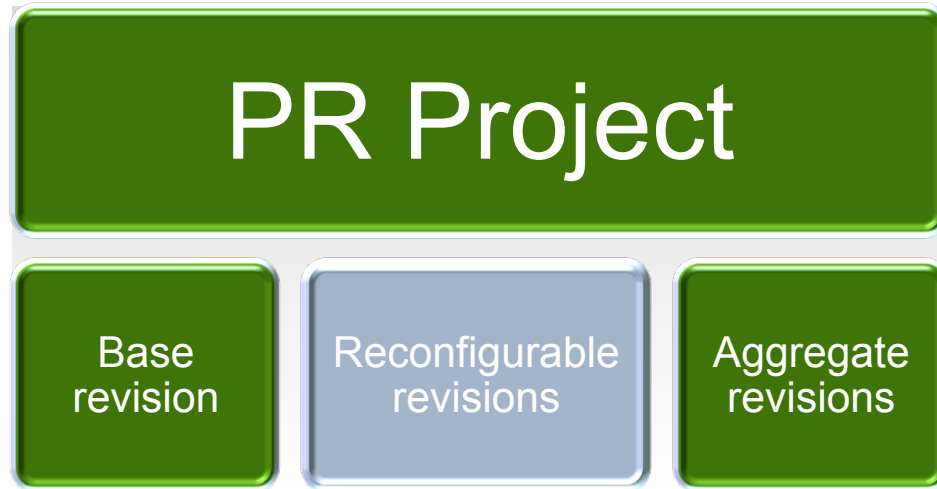
# Partial Reconfiguration Project



## ■ #1: Compile static logic


- User can specify PR1 & PR2 for base revision
- Can be empty, initial content, or “hardest-to-route” personae

# Partial Reconfiguration Project


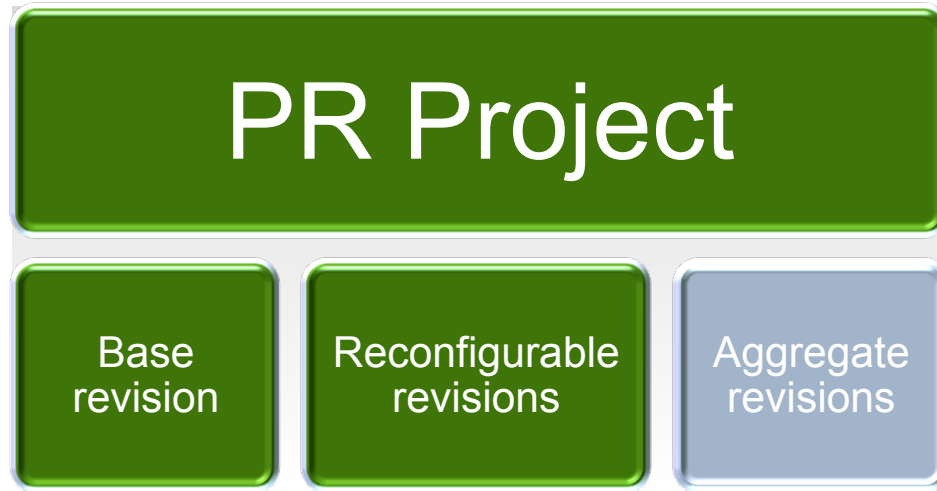


## ■ #2: Compile reconfigurable logic

- Personae compile imports base revision P&R
- Quartus can compile all personae in parallel
  - Combine personae in different regions together
  - # compiles = max # personae across all regions



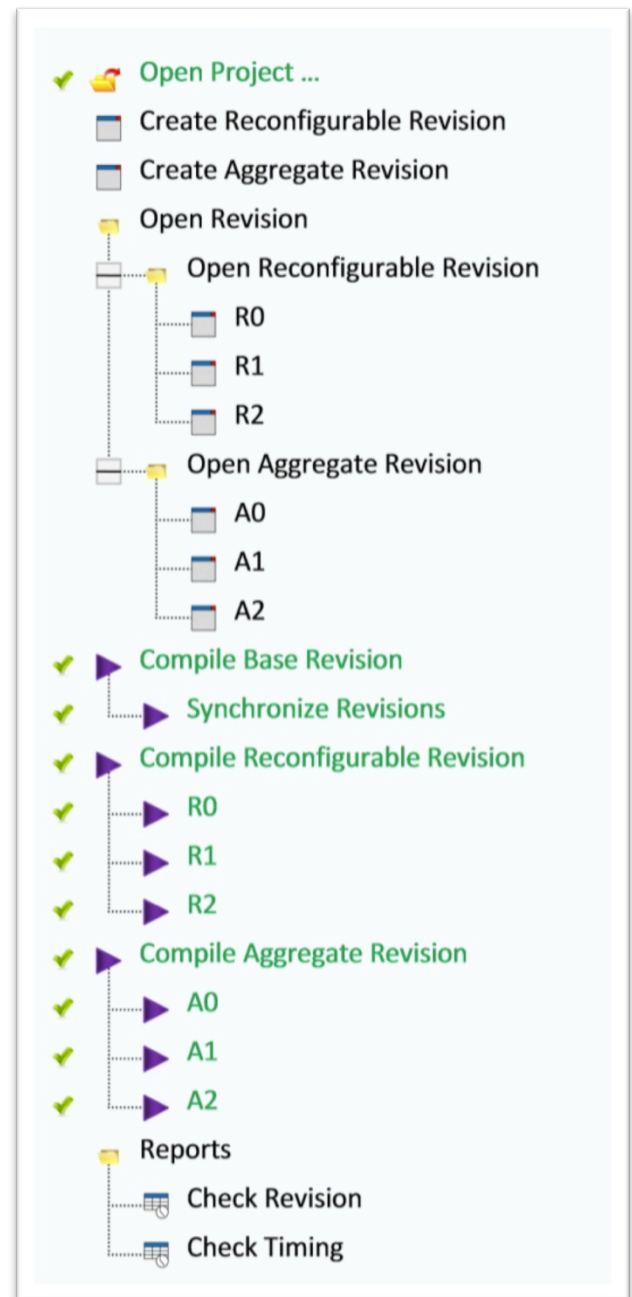
# Partial Reconfiguration Project



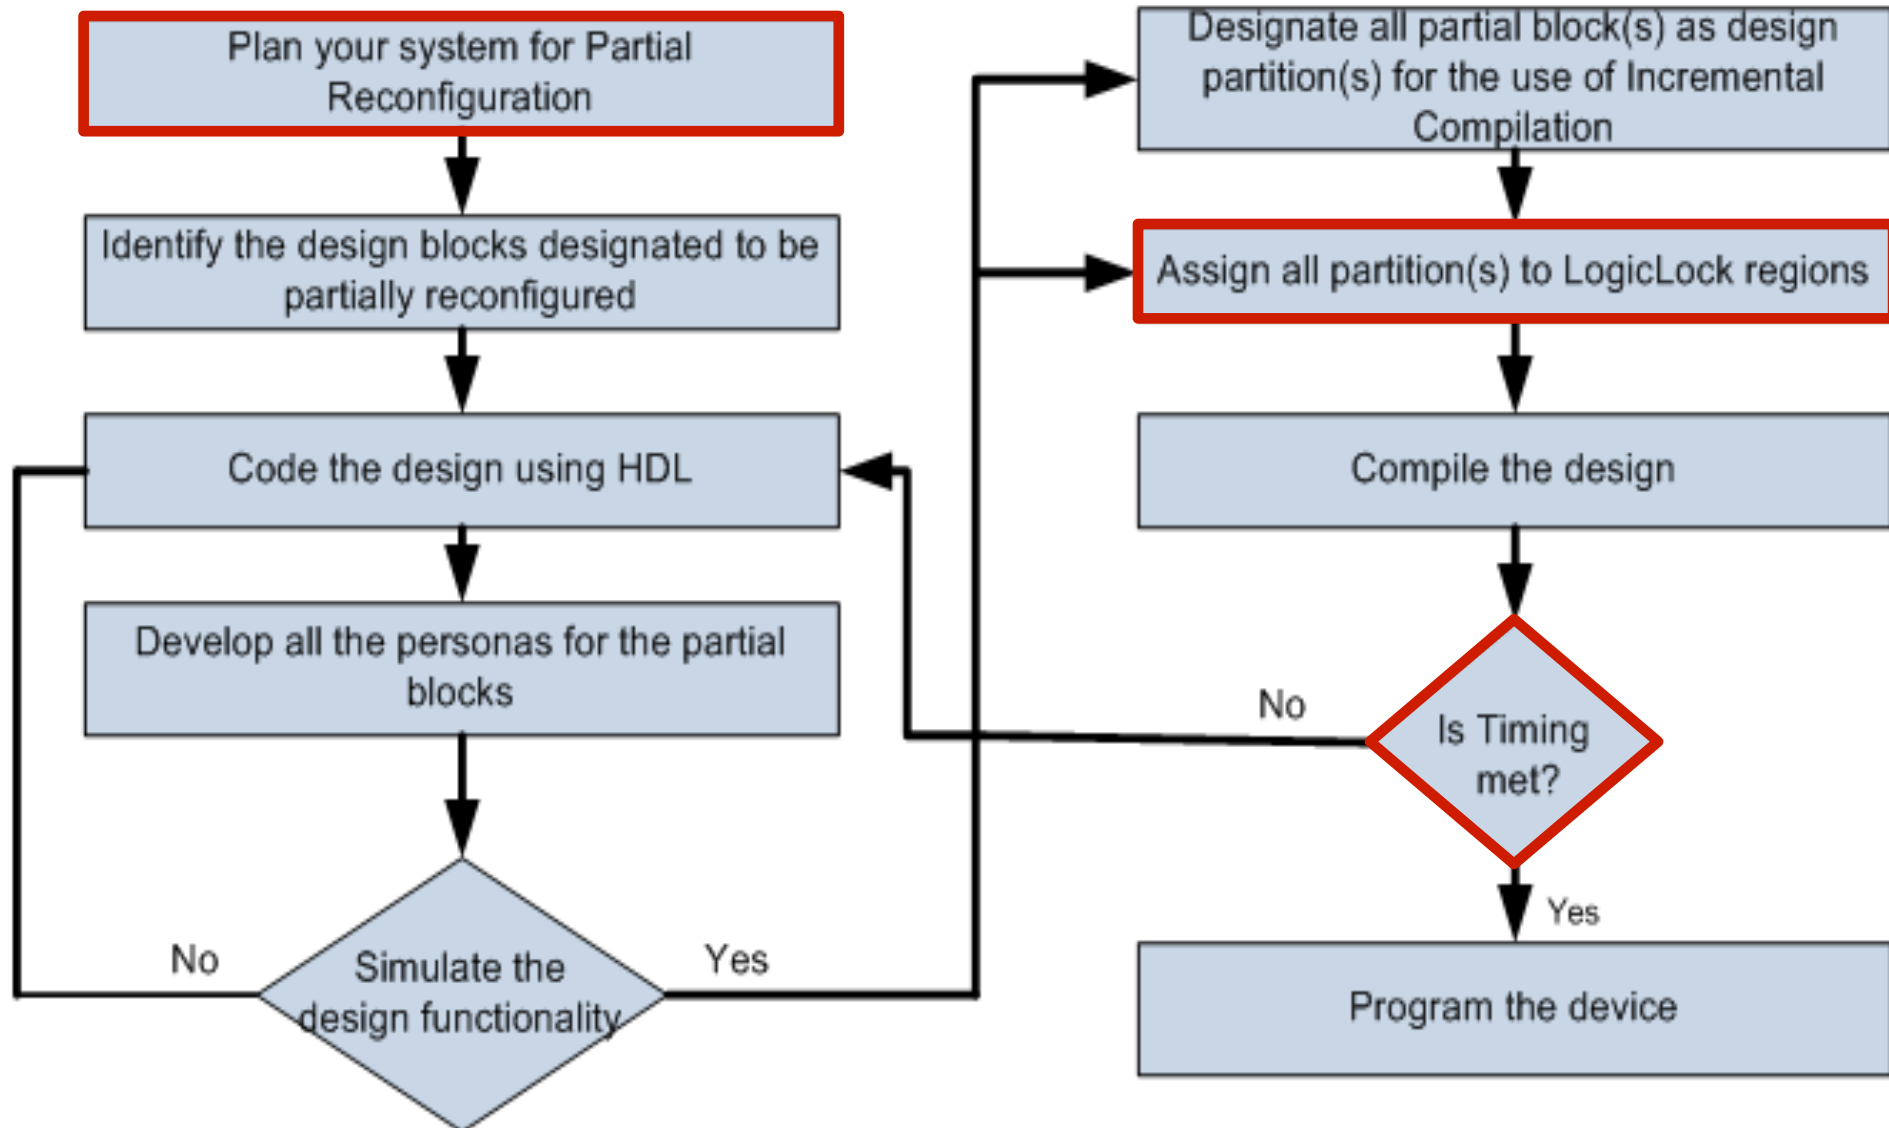
- #3: Compile user-defined groups of reconfigurable logic
  - This step is completely optional
  - Generates complete bitstream, without requiring PR commands
  - Useful for timing analysis / simulation

# Managing Revisions

- Quartus contains dedicated GUI to navigate PR project
  - Can directly launch “regular” Quartus GUI for a specific revision
  - Can easily create new revisions
  - Can “compile all” via single click
  - Can obtain summary of all compiles



# Design Flow




# PR Introduces New Software Challenges

- Handling inputs/output signals to PR partitions
  - PR inputs/outputs must stay in same spot in all compiles
- Partitioning routing resources across all compiles
  - Different PR partitions cannot use same routing resource
  - Different personae in same PR region **can** use same routing wires
- Utilizing shared FPGA resources
  - E.G. pre-fabricated clock networks
- Timing closure on cross-partition paths
  - E.G. reg in PR1 → reg in PR2



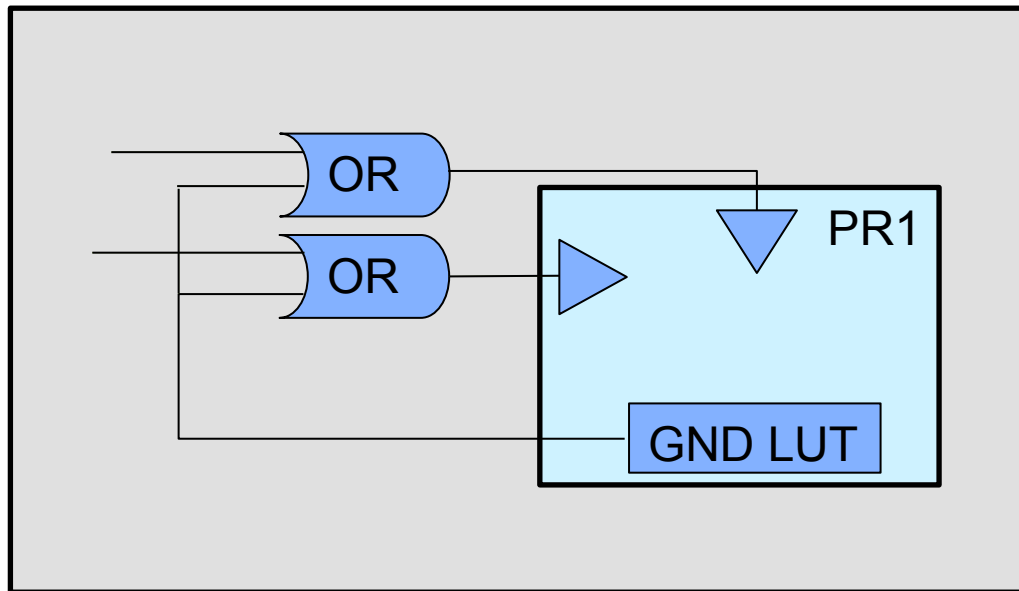
# PR Partition Input Signals

- Recall electrical problem with GND inputs
  - Option 1: User “guarantees” inputs will be VDD when PR engages
  - Option 2: Let compiler solve the problem automatically




# PR Partition Input Signals

- For each PR partition, reserve one LUT for GND
- For each input signal into a PR partition,
  - Add OR gate **outside PR partition** and connect GND LUT to other input
  - Connect OR gate output to a 1-input anchor LUT inside PR partition
- This connectivity is inserted during the base revision



# PR Partition Output Signals

- For each output signal,
  - Create anchor 1-LUT inside PR partition
  - Connect output to logic in static region, according to design files




# Using Anchor 1-LUTs in PR Partitions

- (+) Increases routing flexibility on inputs/outputs
  - FPGA architecture not designed to route to/from specific wires
- (+) P&R is well-defined for each compile
  - Otherwise, how should path from OR gate to logic be chosen?
  - Inputs can branch from the OR gate → 1-LUT path, if desired
- (-) Reduces # of usable LUTs in PR region
- (-) Increases latency on signals entering region

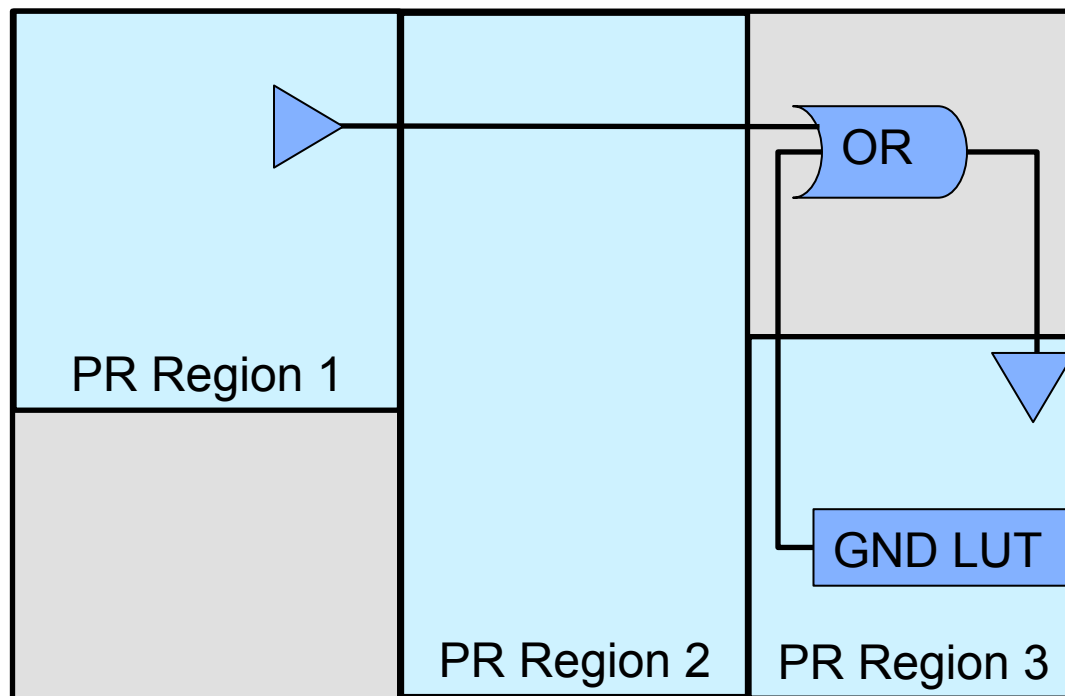
# Route-Throughs At The Top Level

- HW supports reconfiguration of individual routing muxes
  - Enables routing through PR regions
  - Simplifies / removes many floorplanning restrictions
  - Quartus records routing reserved for top-level use and prevents PR instances from using these wires




# Inter-Region Communication

- SW design decisions make this easy to support
  - Anchor LUTs & inserted OR gate are routed in the static compile
  - Resources used by the static compile are disallowed by all personae compilations

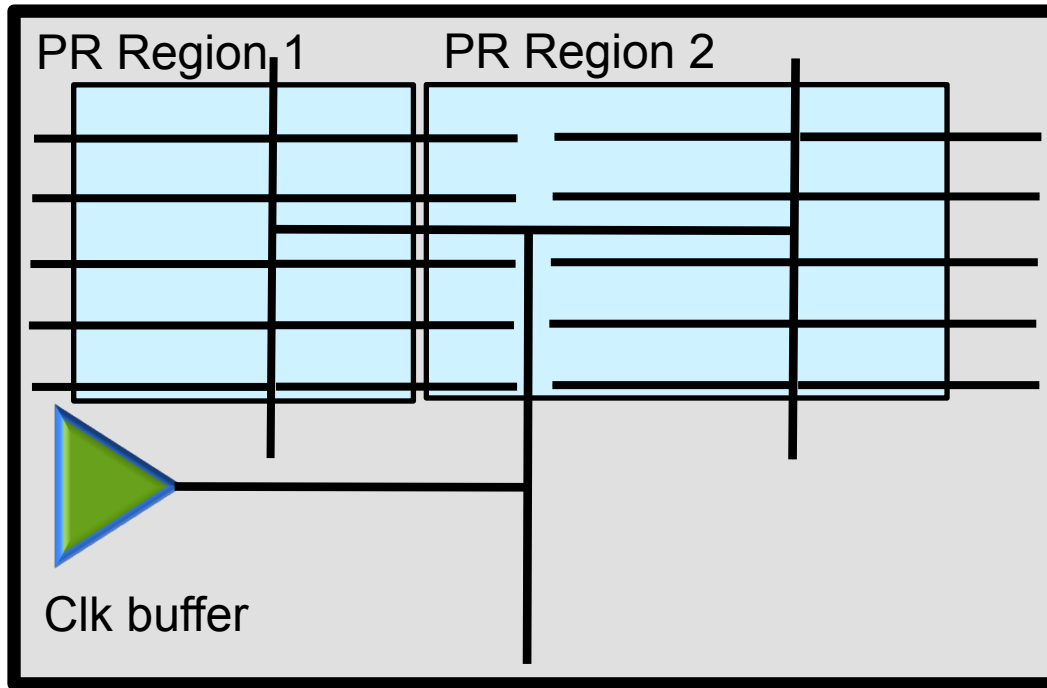


# Global Networks On Stratix V Devices

- 16 global networks
  - Balanced H-tree
  - Reaches every block
  - Only 6 unique clocks can drive each “row clock region” (~12x1)
- PR designs should consider clocking regions
  - Match region size to clock boundaries = easier problem
  - Quartus supports any PR region
  - P&R engine assumes clock feeds every location in PR region

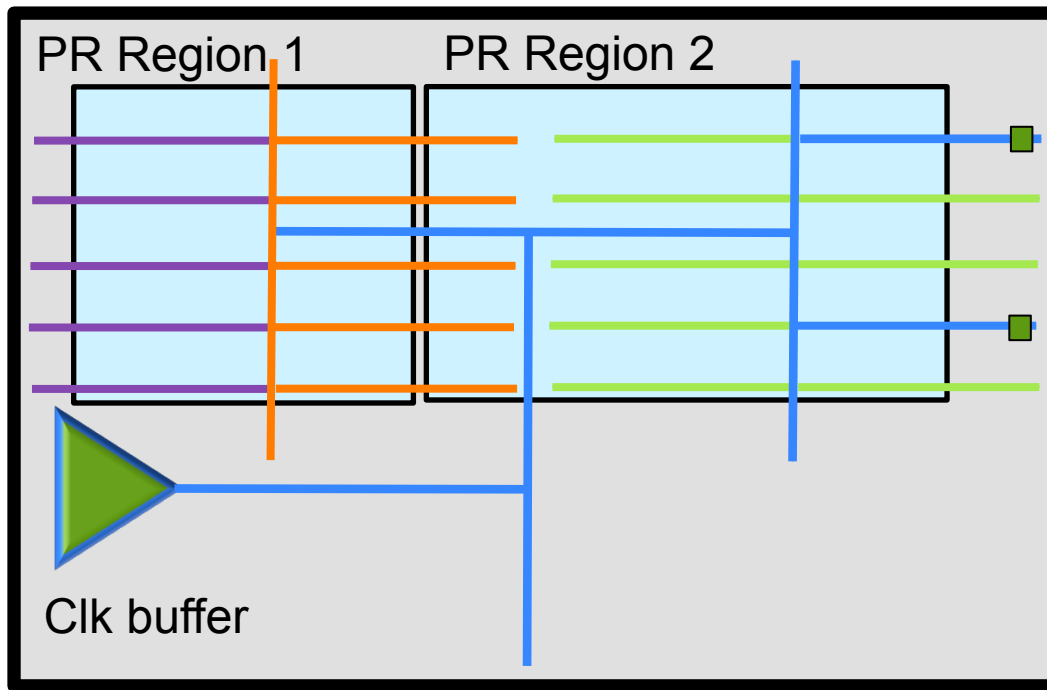


# Sharing Global Wires Between Regions






# Sharing Global Wires Between Regions



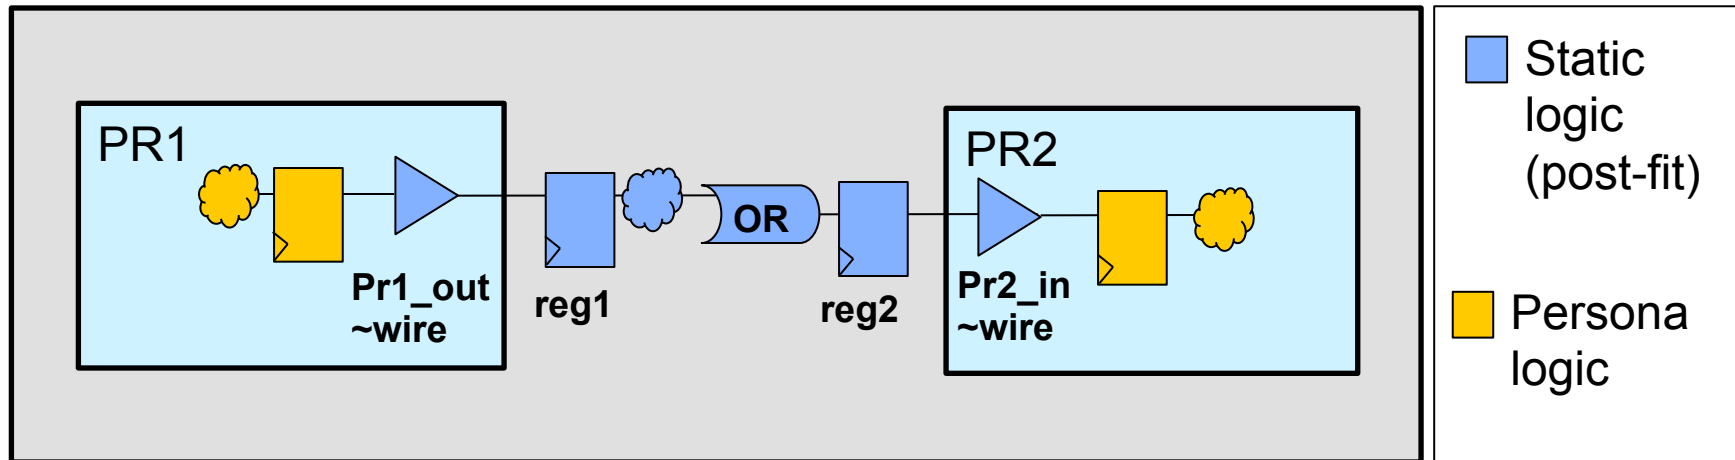
■	Static logic
—	Wires in the base revision bitstream
—	Wires that may be used in a persona bitstream targeting region 2
—	Wires that may be used in a persona bitstream targeting region 1
—	Wires that must be in the base revision bitstream because any persona in region 1 or 2 may use them

# Simple Timing Closure (Base Compile)



- Add registers at all PR partition boundaries
  - Fitter routes intra-region connections for minimum delay
    - pr1\_out~wire → reg1
    - reg2 → pr2\_in~wire
  - 📁 OR gate can be placed before 'reg2' (assuming no clock enable)

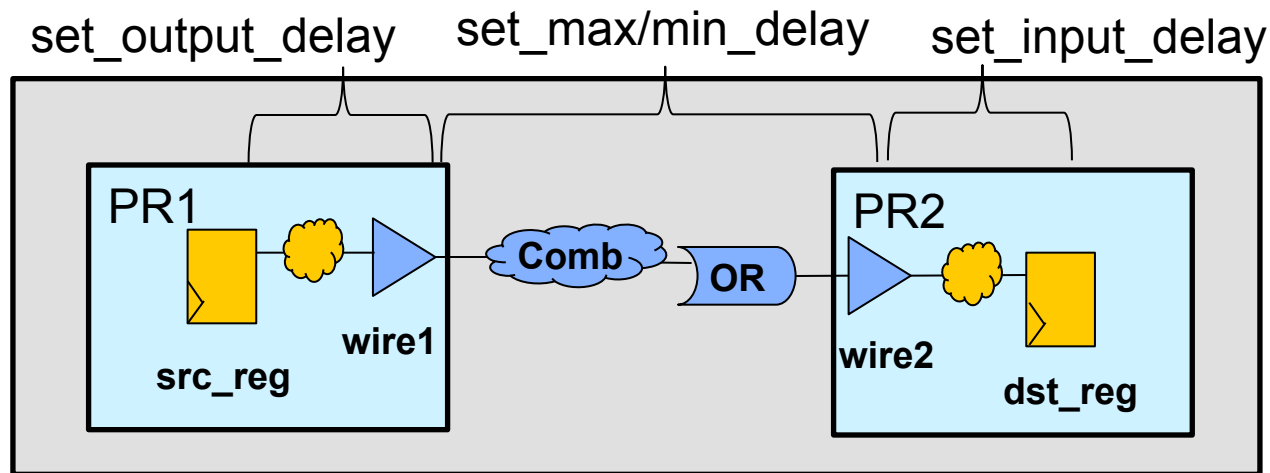
# Simple Timing Closure (Persona Compile)



- Add registers at all PR partition boundaries
  - Compiler sees reg->reg paths and optimizes accordingly
  - 📁 Each path has at least 1 “free” connection in persona compiles
    - Can add more delay, if required, to meet hold timing constraints


# Advanced Timing Closure

- User explicitly manages clock period
  - Divide into three segments
  - Specify timing constraints for each segment in SDC file




# The SDC File

```
set CLK_PERIOD 3.3
set STATIC_BUDGET 2.1
set PR1_BUDGET 0.6
set PR2_BUDGET 0.6
```




The clock period and how the user has divided it

```
create_clock -name {clk} -period $CLK_PERIOD [get_ports {clk}]
create_clock -name {clk_PR1} -period $CLK_PERIOD
create_clock -name {clk_PR2} -period $CLK_PERIOD
```



The clock, and a virtual clock for each input/output domain of the PR partitions

```
set_clock_latency -source -late 2.8 [get_clocks clk_PR1]
set_clock_latency -source -early 2.6 [get_clocks clk_PR1]
set_clock_latency -source -late 2.9 [get_clocks clk_PR2]
set_clock_latency -source -early 2.7 [get_clocks clk_PR2]
```



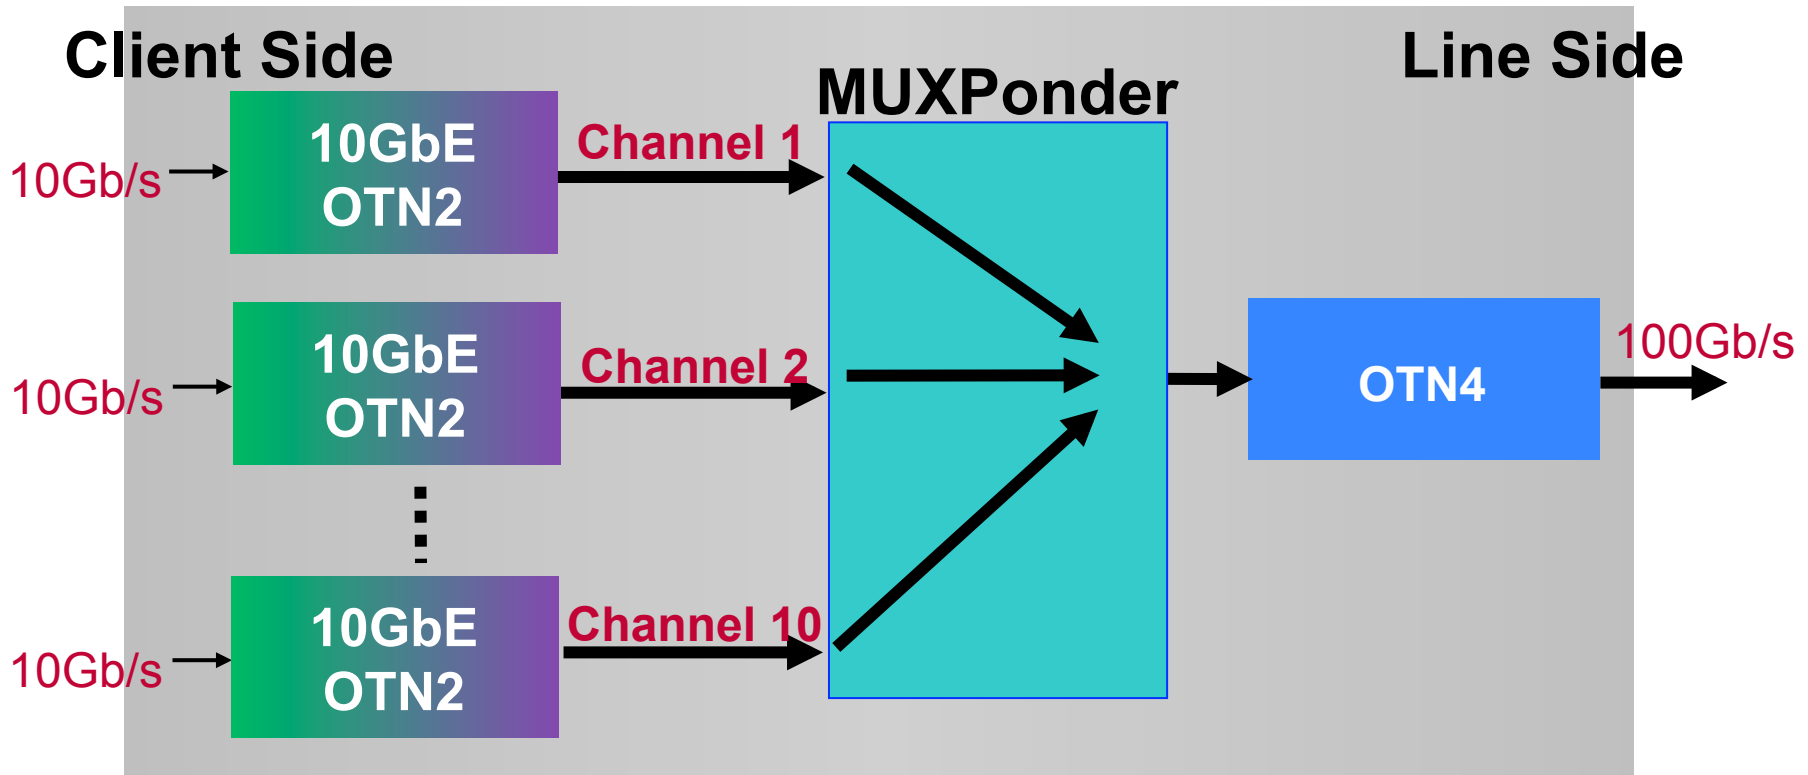
```
set_max_delay -from [get_ports {pr1|out~PR_OPORT}] -to [get_ports {pr2|in1~PR_IPORT}] $STATIC_BUDGET
```

```
set_output_delay -clock clk_PR2 [expr $STATIC_BUDGET + $PR2_BUDGET] [get_ports {pr1|out~PR_OPORT}]
set_input_delay -clock clk_PR1 [expr $STATIC_BUDGET + $PR1_BUDGET] [get_ports {pr2|in1~PR_IPORT}]
```



- User will be provided the worst-case clock arrival times for each global signal entering PR region
  - Will be known at the end of the base revision
  - Generally, these numbers are a function of device/speedgrade


# Upcoming Idea: Translatable Bitstreams



- 10 PR regions, 2 personae per PR region
  - But, there's only 2 unique personae across all regions
  - Users would prefer to generate 2 partial bitstream files, not 20

# Upcoming Idea: Resource Negotiation

- Enforcing hard boundaries of routing resources for PR regions may increase fitting difficulty
  - Guarantees legality at the expense of optimization quality
- Option 1: Partition unused routing resources between PR regions
  - Downside: Increases cost of changing ‘base’
- Option 2: Allow PR regions to use all routing resources and iteratively resolve congestion using costs
  - Downside: Increased latency of overall compile
  - Downside: Significantly increases cost of changing ‘base’



# Upcoming Idea: Minimize Overhead

- Current SW implementation approach works well for multi-modal applications
  - Overhead  $\ll$  PR region size
- Some advanced applications would want to create hundreds of very small regions
  - SW approach doesn't prohibit this design style
  - However, overhead added by SW reduces becomes limiting factor



# Summary

- Altera devices have flexible hardware support for partial reconfiguration
  - Able to reconfigure individual LABs, RAMs, routing muxes
- Altera has provided a well-defined design flow for implementing multi-modal applications using PR
  - Intuitive design entry, simulation, and project cockpit (GUI)
  - User can manage complex timing situations between PR regions
- Quartus automatically handles implementation details efficiently & effectively
  - “OR”-gate insertion
  - Route throughs
  - Cross-region clocking